

In-Lab 9 Report

Optimization Problem

```
#include <iostream>
using namespace std;

int loop(int x, int y){
    int total;
    for(int i=0;i<y;i++){
        total+=x;
    }
    return total;}

int main(){
    int x=5;
    int y=10;
    cout<<loop(x,y)<<endl;
    return 0;
}
```

For the optimization problem I created a small c++ program that runs a loop, adding the parameter x to a total sum y amount of times. I then print out the result in the main method. When looking at the normal code, I can almost completely follow what is happening, however, when looking at the optimized code, I quickly become confused by even the most simple aspects, such as a call. However, I will work through some of this confusion.

Main code optimized on right versus not optimized on the left:

main:	# @main	main:	
	push EBP		push EBP
	mov EBP, ESP		mov EBP, ESP
.Ltmp14:			push EDI
	sub ESP, 24		push ESI
	mov DWORD PTR [EBP - 4], 0		sub ESP, 16
	mov DWORD PTR [EBP - 8], 5	.Ltmp9:	mov DWORD PTR [ESP], _ZSt4cout
	mov DWORD PTR [EBP - 12], 10		call _ZNSolsEi
	mov EAX, DWORD PTR [EBP - 8]		mov ESI, EAX
	mov ECX, DWORD PTR [EBP - 12]		mov EAX, DWORD PTR [ESI]
	mov DWORD PTR [ESP], EAX		mov EAX, DWORD PTR [EAX - 12]
	mov DWORD PTR [ESP + 4], ECX		mov EDI, DWORD PTR [EAX + ESI + 124]
	call _Z4loopii		test EDI, EDI
	lea ECX, DWORD PTR [_ZSt4cout]		je .LBB1_5
	mov DWORD PTR [ESP], ECX	# BB#1:	cmp BYTE PTR [EDI + 28], 0
	mov DWORD PTR [ESP + 4], EAX		je .LBB1_3
	call _ZNSolsEi	# BB#2:	mov AL, BYTE PTR [EDI + 39]
	lea ECX, DWORD PTR		jmp .LBB1_4
[_ZSt4endlcSt11char_traitslcEERSt13basic_ostreamIT_T0_ES		.LBB1_3:	mov DWORD PTR [ESP], EDI
6_]			call _ZNKSt5ctypeclcE13_M_widen_initEv
	mov DWORD PTR [ESP], EAX		mov EAX, DWORD PTR [EDI]
	mov DWORD PTR [ESP + 4], ECX		mov DWORD PTR [ESP], EDI
	call _ZNSolsEPFRSoS_E		mov DWORD PTR [ESP + 4], 10
	mov ECX, 0		call DWORD PTR [EAX + 24]
	mov DWORD PTR [EBP - 16], EAX # 4-	.LBB1_4:	movsx EAX, AL
byte Spill			mov DWORD PTR [ESP + 4], EAX
	mov EAX, ECX		mov DWORD PTR [ESP], ESI
	add ESP, 24		call _ZNSo3putEc
	pop EBP		mov DWORD PTR [ESP], EAX
	ret		call _ZNSo5flushEv
			xor EAX, EAX
			add ESP, 16
			pop ESI
			pop EDI
			pop EBP
			ret

Comparison between main methods:

- The first thing I notice when adding the optimization is that the optimized code makes use of more registers as opposed to offsets of EBP. I am assuming that using more registers allows for faster access.
- The next thing I notice when comparing the two mains is that the call for the loop appears to be missing. After some general research online, I discovered that the loop call is being unwound, which causes the program to still be fairly long (150 lines for non-optimized vs 114 lines for optimized). Wikipedia states, "The goal of loop unwinding is to increase a program's speed by reducing (or eliminating) instructions that control the loop, such as pointer arithmetic and "end of loop" tests on each iteration." This clearly explains the multiple small snippets of x86 code throughout the optimized .s file.

Unoptimized loop left, optimized right:

<pre>_Z4loopii: # @_Z4loopii # BB#0: sub ESP, 16 mov EAX, DWORD PTR [ESP + 24] mov ECX, DWORD PTR [ESP + 20] mov DWORD PTR [ESP + 12], ECX mov DWORD PTR [ESP + 8], EAX mov DWORD PTR [ESP], 0 .LBB1_1: # =>This Inner Loop Header: Depth=1 mov EAX, DWORD PTR [ESP] cmp EAX, DWORD PTR [ESP + 8] jge .LBB1_4 # BB#2: # in Loop: Header=BB1_1 Depth=1 mov EAX, DWORD PTR [ESP + 12] mov ECX, DWORD PTR [ESP + 4] add ECX, EAX mov DWORD PTR [ESP + 4], ECX # BB#3: # in Loop: Header=BB1_1 Depth=1 mov EAX, DWORD PTR [ESP] add EAX, 1 mov DWORD PTR [ESP], EAX jmp .LBB1_1 .LBB1_4: mov EAX, DWORD PTR [ESP + 4] add ESP, 16 ret</pre>	<pre>_Z4loopii: # @_Z4loopii # BB#0: ret From Global: # BB#0: push EBP .Ltmp15: .cfi_def_cfa_offset 8 .Ltmp16: .cfi_offset_ebp, -8 mov EBP, ESP .Ltmp17: .cfi_def_cfa_register ebp sub ESP, 24 mov DWORD PTR [ESP], _ZStL8__ioinit call _ZNSt8ios_base4InitC1Ev mov DWORD PTR [ESP + 8], __dso_handle mov DWORD PTR [ESP + 4], _ZStL8__ioinit mov DWORD PTR [ESP], _ZNSt8ios_base4InitD1Ev call __cxa_atexit add ESP, 24 pop EBP ret</pre>
---	--

Things start to get confusing quickly when comparing the unoptimized loop on the left to the optimized on the right. The first thing I notice is the unoptimized code has four returns and the optimized has only three, probably due to the addition of the loop unwinding. Also, I notice that the loop function (`_Z4Loopii`) is never actually called in the optimized code. Given that it is unnecessary to call it because all it does is run the loop and return. If we are unwinding the loop, then the function no longer

does anything of value, besides returning the total. I also notice that the loop in the unoptimized code creates 16 bytes for local variables. I do not see this sort of allocation for local variables anywhere in the optimized code.

In conclusion, the optimized code, while faster, is extremely difficult to comprehend. Luckily, I was able to find more information about loop unwinding on Wikipedia, which ended up helping to explain some of the strange structure of the assembly code.

Sources:

- http://en.wikipedia.org/wiki/Loop_unwinding
- http://en.wikibooks.org/wiki/X86_Disassembly/Code_Optimization