

In-Lab 9 Report

Optimization Problem

```

#include <iostream>
using namespace std;

int loop(int x, int y){
    int total;
    for(int i=0;i<y;i++){
        total+=x;
    }
    return total;}

int main(){
    int x=5;
    int y=10;
    cout<<loop(x,y)<<endl;
    return 0;
}

```

For the optimization problem I created a small c++ program that runs a loop, adding the parameter x to a total sum y amount of times. I then print out the result in the main method. When looking at the normal code, I can almost completely follow what is happening, however, when looking at the optimized code, I quickly become confused by even the most simple aspects, such as a call. However, I will work through some of this confusion.

Main code optimized on right versus not optimized on the left:

<pre> main: # @main push EBP mov EBP, ESP .Ltmp14: sub ESP, 24 mov DWORD PTR [EBP - 4], 0 mov DWORD PTR [EBP - 8], 5 mov DWORD PTR [EBP - 12], 10 mov EAX, DWORD PTR [EBP - 8] mov ECX, DWORD PTR [EBP - 12] mov DWORD PTR [ESP], EAX mov DWORD PTR [ESP + 4], ECX call _Z4loopii lea ECX, DWORD PTR [_ZSt4cout] mov DWORD PTR [ESP], ECX mov DWORD PTR [ESP + 4], EAX call _ZNSolsEi lea ECX, DWORD PTR [_ZSt4endlcSt11char_traitsIcEERSt13basic_ostreamIT_T0_ES 6_] mov DWORD PTR [ESP], EAX mov DWORD PTR [ESP + 4], ECX call _ZNSolsEPFRSoS_E mov ECX, 0 mov DWORD PTR [EBP - 16], EAX # 4- byte Spill mov EAX, ECX add ESP, 24 pop EBP ret </pre>	<pre> main: push EBP mov EBP, ESP push EDI push ESI sub ESP, 16 .Ltmp9: mov DWORD PTR [ESP], _ZSt4cout call _ZNSolsEi mov ESI, EAX mov EAX, DWORD PTR [ESI] mov EAX, DWORD PTR [EAX - 12] mov EDI, DWORD PTR [EAX + ESI + 124] test EDI, EDI je .LBB1_5 # BB#1: cmp BYTE PTR [EDI + 28], 0 je .LBB1_3 # BB#2: mov AL, BYTE PTR [EDI + 39] jmp .LBB1_4 .LBB1_3: mov DWORD PTR [ESP], EDI call _ZNKSt5ctypeIcE13_M_widen_initEv mov EAX, DWORD PTR [EDI] mov DWORD PTR [ESP], EDI mov DWORD PTR [ESP + 4], 10 call DWORD PTR [EAX + 24] .LBB1_4: movsx EAX, AL mov DWORD PTR [ESP + 4], EAX mov DWORD PTR [ESP], ESI call _ZNSo3putEc mov DWORD PTR [ESP], EAX call _ZNSo5flushEv xor EAX, EAX add ESP, 16 pop ESI pop EDI pop EBP ret </pre>
---	--

Comparison between main methods:

- The first thing I notice when adding the optimization is that the optimized code makes use of more registers as opposed to offsets of EBP. I am assuming that using more registers allows for faster access.
- The next thing I notice when comparing the two mains is that the call for the loop appears to be missing. After some general research online, I discovered that the loop call is being unwound, which causes the program to still be fairly long (150 lines for non-optimized vs 114 lines for optimized). Wikipedia states, "The goal of loop unwinding is to increase a program's speed by reducing (or eliminating) instructions that control the loop, such as pointer arithmetic and "end of loop" tests on each iteration." This clearly explains the multiple small snippets of x86 code throughout the optimized .s file.

Unoptimized loop left, optimized right:

<pre>_Z4loopii: # @_Z4loopii # BB#0: sub ESP, 16 mov EAX, DWORD PTR [ESP + 24] mov ECX, DWORD PTR [ESP + 20] mov DWORD PTR [ESP + 12], ECX mov DWORD PTR [ESP + 8], EAX mov DWORD PTR [ESP], 0 .LBB1_1: # =>This Inner Loop Header: Depth=1 mov EAX, DWORD PTR [ESP] cmp EAX, DWORD PTR [ESP + 8] jge .LBB1_4 # BB#2: # in Loop: Header=BB1_1 Depth=1 mov EAX, DWORD PTR [ESP + 12] mov ECX, DWORD PTR [ESP + 4] add ECX, EAX mov DWORD PTR [ESP + 4], ECX # BB#3: # in Loop: Header=BB1_1 Depth=1 mov EAX, DWORD PTR [ESP] add EAX, 1 mov DWORD PTR [ESP], EAX jmp .LBB1_1 .LBB1_4: mov EAX, DWORD PTR [ESP + 4] add ESP, 16 ret</pre>	<pre>_Z4loopii: # @_Z4loopii # BB#0: ret From Global: # BB#0: push EBP .Ltmp15: .cfi_def_cfa_offset 8 .Ltmp16: .cfi_offset_ebp, -8 mov EBP, ESP .Ltmp17: .cfi_def_cfa_register ebp sub ESP, 24 mov DWORD PTR [ESP], _ZStL8__ioinit call _ZNSt8ios_base4InitC1Ev mov DWORD PTR [ESP + 8], __dso_handle mov DWORD PTR [ESP + 4], _ZStL8__ioinit mov DWORD PTR [ESP], _ZNSt8ios_base4InitD1Ev call __cxa_atexit add ESP, 24 pop EBP ret</pre>
--	--

Things start to get confusing quickly when comparing the unoptimized loop on the left to the optimized on the right. The first thing I notice is the unoptimized code has four returns and the optimized has only three, probably due to the addition of the loop unwinding. Also, I notice that the loop function (`_Z4Loopii`) is never actually called in the optimized code. Given that it is unnecessary to call it because all it does is run the loop and return. If we are unwinding the loop, then the function no longer

does anything of value, besides returning the total. I also notice that the loop in the unoptimized code creates 16 bytes for local variables. I do not see this sort of allocation for local variables anywhere in the optimized code.

In conclusion, the optimized code, while faster, is extremely difficult to comprehend. Luckily, I was able to find more information about loop unwinding on Wikipedia, which ended up helping to explain some of the strange structure of the assembly code.

Sources:

- http://en.wikipedia.org/wiki/Loop_unwinding
- http://en.wikibooks.org/wiki/X86_Disassembly/Code_Optimization

Problem 1: Inheritance

```
class Person{
public:
    Person(string n);
    ~Person();
private:
    string name;
};

Person::Person(string n){
    name=n;
}

Person::~~Person(){
    cout<<"calling ~Person filler code."<<endl;
}

class Teacher: public Person{
public:
    Teacher(string n,int o);
    ~Teacher();
private:
    int office;
};

Teacher::Teacher(string n, int o):Person(n){
    office=o;
}

Teacher::~~Teacher(){
    cout<<"calling ~teacher filler code"<<endl;
}

int main(){
    Teacher t("Bloomfield", 403);
    return 0;
}
```

For problem 1, I wrote a quick c++ program that contains a *Person* and a *Teacher*, similar to the example we went over in class. The person class creates a person object that includes one field, a name. The *Teacher* class extends *Person* and adds an additional field, an office number. The constructor creates a *Person* and then initializes the additional field for *Teacher*. The main method creates a *Teacher* *t* and initializes the values of name to "Bloomfield" and office to "403."

The assembly tends to work in a fairly straightforward manner, following the conventions of c++ inheritance. When I call the main method and create a teacher, main first calls the teacher constructor code, which is copied below.

Main (important portion):

```
lea    EAX, DWORD PTR [EBP - 24]
mov     ECX, ESP
mov     DWORD PTR [ECX + 4], EAX
lea     EAX, DWORD PTR [EBP - 16]
mov     DWORD PTR [ECX], EAX
mov     DWORD PTR [ECX + 8], 403
call    _ZN7TeacherC1ESsi
```

Teacher (important portion):

```
mov  DWORD PTR [EBP - 24], EAX #  
4-byte Spill
```

```
    jmp  .LBB4_2
```

.LBB4_2:

```
    mov  EAX, DWORD PTR [EBP -  
16] # 4-byte Reload  
    mov  DWORD PTR [ESP], EAX  
    call _ZN6PersonD2Ev  
    add  ESP, 56  
    pop  EBP  
    ret
```

The *Teacher* code is copied on the left. Once the registers are correctly adjusted for teacher's additional value, it then calls *Person*, since *Teacher* is just an extension of *Person*.

In terms of memory destruction, the derived class's (*Teacher*) destructor should be called first and then the base class's (*Person*) destructor should be called after, or the reverse order of the constructors (which we verified above). In my code it appears as though there is nothing actually named destructor at the end of the main, where I would expect destruction to occur, however

there is a call to `_ZSt9terminatev`, which I am assuming acts as a destructor. While I cannot locate this in my assembly code, beyond the fact that it is being called, I assume that it follows the destruction methods that I described above.

```
mov  EAX, DWORD PTR [EBP - 20]  
    mov  ECX, DWORD PTR [EBP - 24]  
    mov  DWORD PTR [EBP - 36], ECX #  
4-byte Spill  
    mov  DWORD PTR [EBP - 40], EAX  
# 4-byte Spill  
    jmp  .LBB3_6  
.LBB3_5:  
.Ltmp45:  
    mov  DWORD PTR [EBP - 44], EDX  
# 4-byte Spill  
    mov  DWORD PTR [EBP - 48], EAX  
# 4-byte Spill  
    call _ZSt9terminatev
```

In terms of memory layout, it appears as though the assembly code uses global variables for all of the variables that are static, such as the string "Bloomfield." One interesting memory-related feature I notice in this assembly file is that there are multiple instances where the size of a data type is set using the keyword `size`. For example:

```
.L.str:  
    .asciz  "calling ~Person filler  
code."  
    .size  .L.str, 29
```

I have not seen this before, and I assume that it is related to global constant variables. Also, these global variables never go out of scope, which is why I am assuming they are used in the first place.

Sources for part 2:

<http://stackoverflow.com/questions/120876/c-superclass-constructor-calling-rules>

http://en.wikibooks.org/wiki/X86_Disassembly/Variables#Global_Variables